

# APmap User Manual

Jintao Yu

25th February 2021

## 1 Introduction

APmap is a compiler that compiles a user's applications into the configurations for automata accelerators such as Cache Automaton<sup>1</sup> and RRAM-AP.<sup>2</sup> The compiler targets a two-level routing matrix that consists of global and local switches.

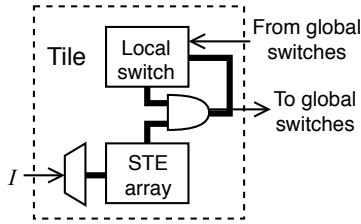


Figure 1: A tile

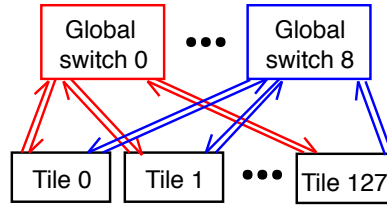


Figure 2: Routing network

The default parameter of APmap models the routing matrix of Cache Automaton. It consists of 128 tiles that each contain a local switch and a state-transit element (STE) array, as shown in Figure 1. Cache Automaton also has eight global switches that each link with all the tiles, as shown in Figure 2.

The network parameters can be configured in APmap via *def.h*. The configurable values include the number of tiles, the number of global switches, the number of STEs in a tile. With this feature, you can generate the configuration fit for your need or explore the hardware design space.

This tool is called using a command

```
./apmap [options] map_file1 [map_file2] ...
```

One or more files are acceptable. The format of the input files are explained in Section 3.1. Three options are available:

- **-h** or **--help**: to show the usage of the tool.
- **--no-g4**: to exclude the 4-way global switch from the routing matrix.

<sup>1</sup>Arun Subramaniyan et al. "Cache Automaton". In: *Proceedings of the 50th Annual IEEE/ACM International Symposium on Microarchitecture*. MICRO-50 '17. Cambridge, Massachusetts: ACM, Oct. 2017, pp. 259–272. URL: <http://doi.acm.org/10.1145/3123939.3123986>.

<sup>2</sup>J. Yu et al. "Memristive devices for computation-in-memory". In: *2018 Design, Automation Test in Europe Conference Exhibition (DATE)*. IEEE, 2018, pp. 1646–1651. DOI: 10.23919/DATE.2018.8342278.

- **--no-opt**: to disable constraint conflict resolving optimizations. This feature is only meant to be used in the evaluation of the tool.

## 2 Build from the source

### 2.1 Dependency

APmap calls METIS<sup>3</sup> to partition the target automata. Therefore, you need to install METIS before building APmap.

### 2.2 Build on Linux

A simple “make” command will build APmap. It has been tested on Ubuntu, CentOS, and Cygwin.

### 2.3 Build on other OS

It should not be challenging to build this project on any mainstream operating system since the source code is written in C.

## 3 Getting started

### 3.1 Input format

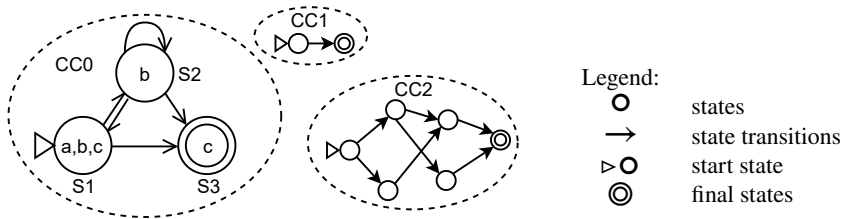


Figure 3: An automaton that consists of three connected components.

The top-level input file is with the suffix *.map*. An automaton can be divided into one or more non-overlapping parts without transitions between any two parts. These parts are named as *connected components* (CCs). For example, Figure 3 contains three CCs, i.e., *CC0*, *CC1*, and *CC2*. It can be described by the following file:

```
1 % automata.map
2 3
3 3 5 cc0.graph
4 2 1 cc1.graph
5 6 7 cc2.graph
```

Any line that starts with a % symbol is comment and will be omitted by APmap. The first line that is not comment specifies the number of CCs in the automaton. From the next line, every line lists three key features of a CC, i.e., the state number, state transition number, and the name of the *.graph* file that describes the details.

<sup>3</sup><http://glaros.dtc.umn.edu/gkhome/metis/metis/overview>

Every line of the *.graph* file describes a state, including the state name, whether it is a start or final state, STE configuration, and the state transitions. The *cc0.graph* mentioned in the last *.map* file represents *CC0* in Figure 3.

```

1 % cc0.graph
2 S1 1 0 00000000 00000000 00000000 E0000000 00000000 00000000 00000000
   00000000 2 3
3 S2 0 0 00000000 00000000 00000000 40000000 00000000 00000000 00000000
   00000000 1 2 3
4 S3 0 1 00000000 00000000 00000000 20000000 00000000 00000000 00000000
   00000000

```

*S1* presented in line 2 is a start state (the ‘1’ after ‘*S1*’), but not a final state (the ‘0’ after ‘1’). The STE configuration field consists of eight eight-digit hexadecimal numbers that are generated based on the input symbols that the state accepts.<sup>4</sup> The last part of a line consists of several line indexes (comment lines are not counted) that the state transits to. In the case of *S1*, it transits to the states described in line index 2 (i.e., *S2*) and line index 3 (i.e., *S3*). If a state has no outgoing transitions such as *S3*, the state transition field is empty.

These files are placed in the *test* folder. You can run this test using the command:

```

$ cd test
$ ../apmap automata.map

```

## 3.2 Output format

After the execution, a file named *map\_result* will be generated in the folder. It describes the hardware configuration of the automata processor. For the example presented in the last section, only a tile is used. The last lines are:

```

281 ...
282 252: S4 0 0 00000000 00000000 00000000 80000000 00000000 00000000
   00000000 00000000 -> 250
283 253: S3 0 0 00000000 00000000 00000000 90000000 00000000 00000000
   00000000 00000000 -> 252
284 254: S2 0 0 00000000 00000000 00000000 A0000000 00000000 00000000
   00000000 00000000 -> 252 251
285 255: S1 1 0 00000000 00000000 00000000 B0000000 00000000 00000000
   00000000 00000000 -> 254 253

```

Each line describes an STE. The first field is the index, which is between 0 and 255. The following fields are almost the same as those in the *.graph* files, except that the destination STEs are referred to using STE indexes.

The configuration of global switches is similar. Nevertheless, it is not shown in this example, as no global switches are used.

## 3.3 Toolchain integration

For large automata, it would be difficult for composing the input files manually. Therefore, we recommend a stack flow with other tools. RAPID<sup>5</sup> is a high-level programming

<sup>4</sup>P. Dlugosch et al. “An Efficient and Scalable Semiconductor Architecture for Parallel Automata Processing”. In: *IEEE Transactions on Parallel and Distributed Systems* 25.12 (2014), pp. 3088–3098. ISSN: 1045-9219. DOI: 10.1109/TPDS.2014.8.

<sup>5</sup><https://github.com/kevinaangstadt/rapid>

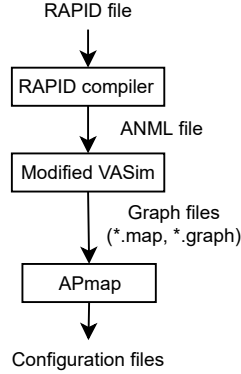


Figure 4: Recommended toolchain.

language developed for automata processing. It can be used to describe the application, and its compiler generates ANML files as the output. VASim<sup>6</sup> can parse ANML files and conduct optimizations. We modified VASim to provide the input files for APmap. The toolchain is shown in Figure 4.

The modified VASim files are placed in the *vasim* folder, i.e., *main.cpp*. We add an option “--apmap” (short form: “-A”) for VASim to generate the *.map* and *.graph* files for APmap.

This modification is based on the latest version of VASim (version 4c615cb). You can download the source code of VASim, replace *main.cpp* with the one we provide and compile. To convert automata to the APmap format, you can use the following command:

```
$ vasim -A xxx.anml
```

Other options, such as “-O”, can also be applied together with “-A”.

## 4 Source Insight

The general workflow of APmap is shown in Figure 5. APmap first reads the state numbers of all the CCs from the *.map* file and sort them in descending order. Then, APmap picks the largest unmapped CC and maps it to one or multiple tiles. In some cases, not all the space of these tiles are occupied by this CC. Therefore, APmap tries to find some small CCs to fill in the remaining space. This process repeats until all the CCs are mapped.

When a CC contains more than 256 states (i.e., the size of an STE array), it has to be mapped to multiple tiles, following the workflow shown in Figure 6. First, this CC is partitioned into several parts. To increase the chance of mapping success, the partitioning process produces multiple solutions. Then, APmap examines the partitioning results with the hardware constraints on a tile. If there are conflicts, an extra process is applied to resolve the conflicts. Next, APmap tries to generate the configuration of the

<sup>6</sup><https://github.com/jackwadden/VASim>

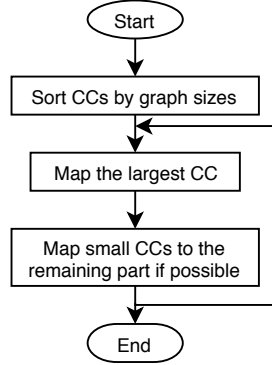


Figure 5: APmap workflow.

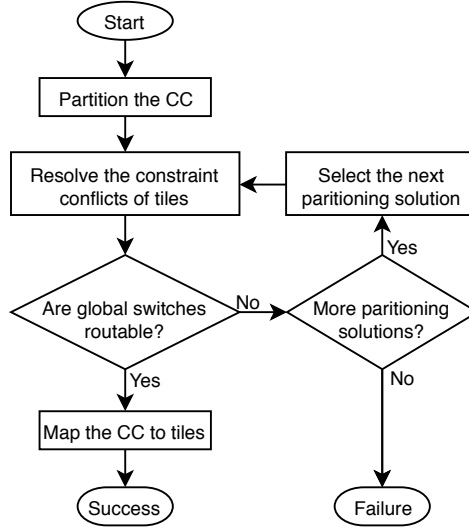


Figure 6: Workflow of mapping a CC that contains more than 256 states.

global switches that connect the CC's partitioned parts. If it succeeds, the configuration of tiles will subsequently be generated. In this case, a valid configuration for the target CC has been produced. If the configuration of global switches cannot be generated, APmap selects the next partitioning solution and repeats the previous processes. If none of these partitioning solutions leads to a valid configuration, the mapping flow fails.

When a CC is larger than the capacity of a tile, it needs to be partitioned into several parts. APmap conducts this process with the help of METIS. METIS divides an undirected graph into  $k$  non-overlapping parts while trying to cut the least number of edges. An edge is *cut* in a division means that the two nodes linked by the edge are assigned to different parts. Here,  $k$  and the number of nodes contained in each part, i.e., the *size* of each part, are specified by the user. If an edge is cut during the partitioning, it will be mapped to a global switch. Less edges are cut generally means that less switching resources are required for the mapping. Therefore, the optimizing target of METIS is

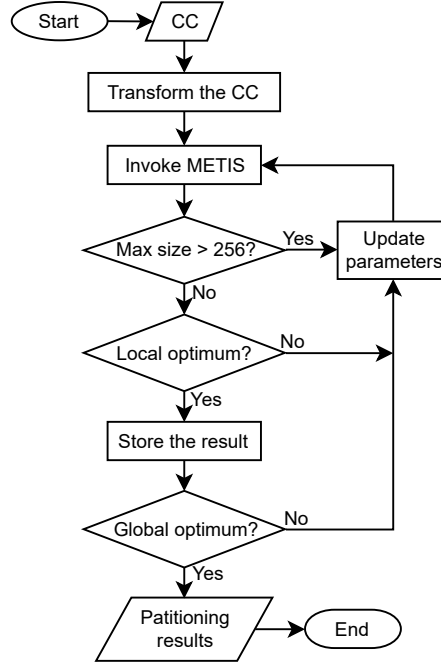


Figure 7: Workflow of partitioning a CC.

aligned with our need.

The workflow of the partitioning process is illustrated in Figure 7. First, we transform the CC to make it acceptable for METIS. Then, we invoke METIS to do the actual partitioning. This process is iterated with different parameters to acquire optimal partitioning results that fit our needs. Finally, these partitioning results are returned for further processing.

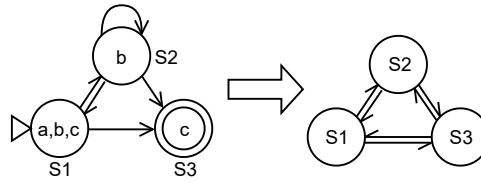


Figure 8: Transforming a CC into an undirected graph.

METIS only works with undirected graphs; therefore, the CC needs to be transformed before partitioning. The transformation contains three major steps. First, APmap constructs an undirected graph based on the transition function of the CC. It is straightforward as we only need to create a node and an edge for every state and transition, respectively. CC's other properties, such as the input symbols associated with the states, the start and final states, are omitted. Next, self-loops are removed as they are not supported by METIS. Finally, if there is only one edge between two nodes, another edge is added in the opposite direction. Figure 8 uses *CC0* in Figure 3 as an example to illustrate this transformation. Note that the self-loop on *S2* has been removed, and two additional

edges (i.e., from  $S3$  to  $S2$  and from  $S3$  to  $S1$ ) are added. There were already a pair of edges between  $S1$  and  $S2$ ; therefore, they remain unchanged.

After the undirected graph is generated, METIS is invoked with two types of input parameters: the number of parts and the size constraints on those parts. The size of a part means the number of nodes contained in that part, and a size constraint is the desired size of a part. METIS produces the same number of parts as the input parameter; however, the size constraints are not guaranteed to be satisfied. It is designed this way because forcing the results to meet the size constraints may lead to a very bad partitioning, i.e., too many edges being cut. Instead, METIS tries to balance constraint satisfaction and partitioning quality. Therefore, we need to check the size of each part after the partitioning. If any part contains more than 256 nodes, it cannot be mapped to a tile. In this case, we need to modify parameters and invoke METIS again.